

NCI Australia's Jamie's GPU Hackathon

A lightning talk

Jamie Border

August 26th, 2020

Why so slow?

Num Calls	Time Time	Func Time	Per Cat		Filter out	% of total
1	30102	4	4	_Domain		0.013373901
1	30095	0	0	simulate.Simulation.simulate.Simulation_ctor[immutable(char)]		0
1	30075	28	28	void simulate.Simulation.run()		0.093617306
261	13546	37	0	void simulate.Simulation.applyWFRPolynomialF()		0.122708563
16905	13135	1210	0	void reconstruction.WFRScheme.WalshFunctionPCFOverElement[finiteVolume.cell.FVCelWFR[], finiteVolume.face.FVInterfaceWFR[], globalConfig.GlobalConfig[], const(bool), finiteVolume		4.040565002
522	10516	1178	2	void plot.plotter.generatePlotData(double[], double[][], immutable(char)[], immutable(char[]), bool, immutable(char[])		3.93613795
261	6738	103	0	void simulate.Simulation.writeDataOut[("gas.rho", "vel", "gas.p", "gas.u"), writeDataOut(ulong, bool)		0.344377946
32219331	5790	1094	0	pure nothrow @nogc @trusted double std::math::pow(double, double) pow(double, double)		3.657701878
67630	5070	1545	0	double[] Walsh.FWT.FWT_x[ulong, ulong, double, double, double]		5.165669197
32219331	4698	3697	0	pure nothrow @nogc @safe real std::math::pow(double, double) pow(double, double) impl(real, real)		12.36082784
17487	4499	73	0	double[] Walsh.PCF.solveCoeffsWithShockDetectionFromPoints(ulong, ulong, double, double, double[], ref bool, ref double, ref double, ref bool, double)		0.24407369
268840	4135	3201	0	@safe void std::file.append(immutable(char[]), append(immutable(char[]), const(void[]))		10.76246414
261	3095	64	0	void simulate.Simulation.calculateEinfelds(const(ulong), body)		0.21392413
33408	3169	4	0	double[] Walsh.PCF.generatePolyCoeffsFromPoints(ulong, ulong, double, double, double[], double[])		0.013373901
260	2850	2	0	void simulate.Simulation.calculateInvscdFluxesFromTimeLevelWFR(const(ulong))		0.0066895
16900	2661	14	0	void fluxCalculators.WFRFluxCalc.calculateFluxWFR[finiteVolume.cell.FVCelWFR[], finiteVolume.face.FVInterfaceWFR[], const(ulong), const(bool)]		0.046809653
137540	2239	67	0	void finiteVolume.face.FVInterfaceWFR.calculateInvscdFluxWFR(Corrections)		0.224012839
260	1317	18	0	void simulate.Simulation.reconstructFromTimeLevelWFR(const(ulong))		0.060182554
85007	1056	359	0	double[] Walsh.FWT.FWT_BTofAFromPoints(const(ulong), const(double), const(double), const(double))		1.2003076
140352	995	501	0	@trusted double Walsh.corr.fg_n(const(ulong), const(double), const(double), const(double))		1.675081079
771692110	000	000	0	===== nothrow @nogc @safe real std::math::pow(double, double) pow(double, double) =====		===== nothrow


$$f_i = \sum_{n=1}^{2^p} A_n g_n(x_i)$$

$$x_i^m = \sum_{n=1}^{2^p} B_{n,m} g_n(x_i)$$

with $1 \leq i \leq 2^p$, $m \geq 0$ and where

$$A_n = \sum_{i=1}^{2^p} g_n(x_i) f_i \Delta x$$

$$B_{n,m} = \sum_{i=1}^{2^p} g_n(x_i) x_i^m \Delta x.$$




Thread



Block



Grid



CUDA C Kernel



```

if (tid < N) {
    F1 = fi[(tid / 32) * 32 + seqi];
}

int Nm = Na/2;
for(int pm=0;pm<Pa;pm++) {
    // calculate negMask
    negMask = (((tid >> (Pa-pm-1)) & 1LU) * 2 - 1) * -1; // 1 or -1

    // calculate src mask
    srcMask = ((tid >> (Pa-pm-1)) & 1LU) ^ 1LU; // 0 or 1

    // apply warp shuffle down
    F2 = srcMask * __shfl_down_sync(0xFFFFFFFF, F1, Nm);

    // flip mask
    srcMask ^= 1LU;

    // apply warp shuffle up
    F2 += srcMask * __shfl_up_sync(0xFFFFFFFF, F1, Nm);

    // add to existing warp value, using negMask
    F1 = F1 * negMask + F2;

    // update shfl width
    Nm >>= 1;

    // write to global memory
    if (tid < N) {
        Fa[tid] = F1;
    }
}

```

Results

